

Comparison Of Two Tests

Zheng Yuan

2019/8/12

```
library(simstudy)
library(dplyr)
library(caret)
library(lmtest)
library(knitr)
```

Define leave one out cross validation prediction error function

Leave one out cross validation prediction error is evaluated by `loocv()` function, which is built based on

$$LOOCV = \frac{1}{n} \sum_{i=1}^n \frac{(Y_i - \hat{Y}_i)^2}{(1 - h_{ii})^2}$$

```
#### leave one out cross validation prediction error function using trick

loocv=function(fit){

  h=lm.influence(fit)$h
  mean((residuals(fit)/(1-h))^2)

}
```

Define likelihood ratio test function

(1)First step: Evaluate the maximum likelihood estimation of σ^2 , that is

$$\hat{\sigma}_\alpha^2 = \frac{1}{n} RSS_\alpha = \frac{1}{n} (Y - X_\alpha \hat{\beta}_\alpha)^T (Y - X_\alpha \hat{\beta}_\alpha)$$

for any certain model M_α . Let's call this function '`sd_mle()`'.

(2)Second step: Assuming normality, the maximum log-likelihood value for model α , $l(\hat{\beta}_\alpha, \hat{\sigma}_\alpha^2)$ say, for the full-rank regression model is the log probability density of Y with maximum likelihood estimation as parameters, namely,

$$\begin{aligned} l(\hat{\beta}_\alpha, \hat{\sigma}_\alpha^2) &= \log \left\{ (2\pi\hat{\sigma}_\alpha^2)^{-\frac{n}{2}} \exp \left\{ -\frac{1}{2\hat{\sigma}_\alpha^2} \|Y - X_\alpha \hat{\beta}_\alpha\|^2 \right\} \right\} \\ &= -\frac{n}{2} \log(2\pi\hat{\sigma}_\alpha^2) - \frac{1}{2\hat{\sigma}_\alpha^2} \|Y - X_\alpha \hat{\beta}_\alpha\|^2 \\ &= \sum_{i=1}^n -\frac{1}{2} \log(2\pi\hat{\sigma}_\alpha^2) - \frac{1}{2\hat{\sigma}_\alpha^2} (Y_i - X_{i,\alpha} \hat{\beta}_\alpha)^2 \\ &= \sum_{i=1}^n dnorm(Y_i; X_{i,\alpha} \hat{\beta}_\alpha, \hat{\sigma}_\alpha^2) \end{aligned}$$

, where $l(.,.)$ is the log-likelihood function.

Let's call this function 'ml()'.

(3) Third step: Evaluate log likelihood ratio test statistic according to definition, namely,

$$LLR = -2(l(\hat{\beta}_\alpha, \hat{\sigma}_\alpha^2) - l(\hat{\beta}_\lambda, \hat{\sigma}_\lambda^2))$$

, where we assume $M_\alpha \subset M_\lambda$. Let's call this function 'lrt()'.

```
sd_mle <- function(object) {  
  sqrt(mean(residuals(object)^2))  
}  
  
ml <- function(object) {  
  sum(dnorm(model.response(model.frame(object)),  
    mean = fitted(object), sd = sd_mle(object), log = TRUE))  
}  
  
lrt <- function(model1,model2) {  
  -2 * (ml(model1) - ml(model2))  
}
```

Define A Full Rank Dataset

The true linear model is $y=3+2x_1+5x_2+3x_3+1.5x_4+e$, which is the full model.

```
#### Define A Full Rank Dataset  
  
def <- defData(varname = "x1", dist="uniform",formula = "10;20") ## x1 is from unifrom distribution  
  
def <- defData(def,varname = "x2", dist="uniform",formula = "0;3")  
  
def <- defData(def,varname = "x3", dist="uniform",formula = "0;5")  
  
def <- defData(def,varname = "x4", dist="uniform",formula = "5;10")  
  
def <- defData(def, varname = "y", formula = "3+2*x1+5*x2+3*x3+1.5*x4", variance = 1)  
  
## The true linear model is y=3+2*x1+5*x2+3*x3+1.5*x4+e
```

Generate a dataset with size 1000

Here I wanted to do a single time simulation to see exactly what values the test statistics are under the alternative hypothesis, i.e., under full model. In this simulation, I generated a dataset of size 1000 and fit model 1~4, respectively.

To test the power of a certain test, basically we set the null hypothesis as model 1~3 respectively, because they are all incorrect (insufficient) model compared with the model under alternative hypothesis, which we set as the true model. The goal is to see exactly to what extent can we reject the null hypothesis in each test.

```
md<-c()  
  
n = 1000
```

```

dt <- genData(n, def)
dt <- dt%>%select(y,x1,x2,x3,x4) ##generate dataset n=1000
fit1 <- lm(y~ x1, data = dt)
fit2 <- lm(y ~ x1+x2, data = dt)
fit3 <- lm(y ~ x1+x2+x3, data = dt)
fit4 <- lm(y ~ x1+x2+x3+x4, data = dt)
md<-c(loocv(fit1),loocv(fit2),loocv(fit3),loocv(fit4)) ##loocv prediction error for each model

```

Calculate the test statistic in one time simulation

Evaluate each test statistic based on

$$\text{New Test Statistic} = \left(\frac{n}{\sigma^2}\hat{\Gamma}_{\alpha,n} - 2d_{\alpha}\right) - \left(\frac{n}{\sigma^2}\hat{\Gamma}_{\lambda,n} - 2d_{\lambda}\right), \text{ and}$$

$$LLR = -2(l(\hat{\beta}_{\alpha}, \hat{\sigma}_{\alpha}^2) - l(\hat{\beta}_{\lambda}, \hat{\sigma}_{\lambda}^2))$$

and compare the corresponding value of two test statistics.

```

ts1 = (n*md[1]-2*1) - (n*md[4]-2*4)
ts2 = (n*md[2]-2*2) - (n*md[4]-2*4)
ts3 = (n*md[3]-2*3) - (n*md[4]-2*4)

lr1 = lrt(fit1,fit4)
lr2 = lrt(fit2,fit4)
lr3 = lrt(fit3,fit4)

ts1;lr1;ts2;lr2;ts3;lr3

## [1] 40661.32
## [1] 3759.252
## [1] 24493.39
## [1] 3265.798
## [1] 4579.586
## [1] 1740.793

```

Now do the same simulation once again, only with sample size of 5000 this time. ## Generate a dataset with size 5000

```

md<-c()

n = 5000

```

```

dt <- genData(n, def)

dt <- dt%>%select(y,x1,x2,x3,x4) ##generate dataset n=1000

fit1 <- lm(y~ x1, data = dt)

fit2 <- lm(y ~ x1+x2, data = dt)

fit3 <- lm(y ~ x1+x2+x3, data = dt)

fit4 <- lm(y ~ x1+x2+x3+x4, data = dt)

md<-c(loocv(fit1),loocv(fit2),loocv(fit3),loocv(fit4)) ##loocv prediction error for each model

```

Calculate the test statistic in one time simulation

```

ts1 = (n*md[1]-2*1) - (n*md[4]-2*4)

ts2 = (n*md[2]-2*2) - (n*md[4]-2*4)

ts3 = (n*md[3]-2*3) - (n*md[4]-2*4)

lr1 = lrt(fit1,fit4)

lr2 = lrt(fit2,fit4)

lr3 = lrt(fit3,fit4)

ts1;lr1;ts2;lr2;ts3;lr3

```

```

## [1] 212415.4
## [1] 18906.95
## [1] 117538
## [1] 16037.23
## [1] 23609.33
## [1] 8756.283

```

Interpretation: from the test statistics value above, we can see that in this one-time simulation, our new test statistics are much greater than their corresponding traditional maximum likelihood ratio test statistics, which implies

that our new test enjoys a much higher power. Besides, the difference becomes larger as the difference between dimensions and sample size n become larger.

1000 times simulations

Now let's make 1000 times simulations, in each time to see whether our new test statistic is greater than its counterpart log likelihood ratio test statistic. Since they share the same criterion, a larger value implies a greater power of the test.

```

df1 <- c()
df2 <- c()
df3 <- c()

for(i in 1:1000){

  dt <- genData(100, def)%>%select(y,x1,x2,x3,x4)
  ##generate dataset n=100
  ##The true linear model is  $y=3+2*x1+5*x2+3*x3+1.5*x4+e$ 
  fit1 <- lm(y~ x1, data = dt)

  fit2 <- lm(y ~ x1+x2, data = dt)

  fit3 <- lm(y ~ x1+x2+x3, data = dt)

  fit4 <- lm(y ~ x1+x2+x3+x4, data = dt)

  md<-c(loocv(fit1),loocv(fit2),loocv(fit3),loocv(fit4))

  ts1 = (n*md[1]-2*1) - (n*md[4]-2*4) ##new test statistics
  ts2 = (n*md[2]-2*2) - (n*md[4]-2*4)
  ts3 = (n*md[3]-2*3) - (n*md[4]-2*4)


  lr1 = lrt(fit1,fit4)
  lr2 = lrt(fit2,fit4)
  lr3 = lrt(fit3,fit4)

  df1 <- c(df1,ifelse (ts1 > lr1, 1, 0 ) )
  df2 <- c(df2,ifelse (ts2 > lr2, 1, 0 ) )
  df3 <- c(df3,ifelse (ts3 > lr3, 1, 0 ) )
  ## In each simulation record if new test statistics is larger than traditional maximum likelihood ratio
}

```

Show the result

```
table(df1)
```

```
## df1
##      1
## 1000
```

```
table(df2)
```

```
## df2
```

```
##      1  
## 1000  
table(df3)
```

```
## df3  
##      1  
## 1000
```

Interpretation: from the result, we can see that in this 1000 times mulations, our new test statistics are greater than traditional maximum likelihood ratio test statistics each single time, which implies that the new test has the same level as maximum likelihood ratio tes while enjoys a higher power and its performance is very stable. So it can beat likelihood ratio test on power.